

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.preprocessing import LabelEncoder
from google.colab import files
import warnings
warnings.filterwarnings('ignore')

# Upload CSV file
uploaded = files.upload()
path = list(uploaded.keys())[0]
df = pd.read_csv(path, encoding_errors="ignore")

print(f"Original number of rows: {len(df)}")
print("Original columns:", df.columns.tolist())

# Drop rows with missing key values
df = df.dropna(subset=['price', 'sqft_living', 'bedrooms', 'bathrooms', 'yr_built', 'date'])

df['date'] = pd.to_datetime(df['date'], errors='coerce')
df = df.dropna(subset=['date'])

# Extract year and month, extract zip code
df['year'] = df['date'].dt.year
df['month'] = df['date'].dt.month
df['zip'] = df['statezip'].astype(str).str.split().str[-1]

# Filter unrealistic values
df = df[df['bedrooms'] > 0]
df = df[df['bathrooms'] > 0]
df = df[(df['sqft_living'] > 100) & (df['sqft_lot'] > 100)]
df = df[(df['floors'] > 0) & (df['floors'] <= 4)]
df = df[df['yr_built'] >= 1900]
df = df[df['price'] > 0]

# Function to remove outliers using IQR
def remove_outliers(df, col):
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
    return df[(df[col] >= lower) & (df[col] <= upper)]

df = remove_outliers(df, 'price')
df = remove_outliers(df, 'sqft_living')
df = remove_outliers(df, 'sqft_lot')

print("\nNumber of rows after cleaning:", len(df))
print(df.describe[['price', 'sqft_living', 'bedrooms', 'bathrooms']])

# Feature engineering
df['age'] = 2014 - df['yr_built']
df['is_renovated'] = (df['yr_renovated'] > 0).astype(int)
df['total_rooms'] = df['bedrooms'] + df['bathrooms']
df['living_to_lot_ratio'] = df['sqft_living'] / df['sqft_lot']

le_city = LabelEncoder()
df['city_encoded'] = le_city.fit_transform(df['city'])

core_features = [
    'sqft_living', 'bathrooms', 'bedrooms', 'sqft_lot',
    'floors', 'waterfront', 'view', 'condition',
    'age', 'is_renovated', 'total_rooms', 'living_to_lot_ratio', 'city_encoded'
]

X = df[core_features]
y = df['price']

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
print(f"\nTraining set: {len(X_train)} records, Test set: {len(X_test)} records")

# Build Random Forest model
rf_model = RandomForestRegressor(
    n_estimators=100,
    max_depth=10
)

```

```

max_depth=10,
min_samples_split=20,
min_samples_leaf=10,
max_features=0.5,
random_state=42,
n_jobs=-1
)

rf_model.fit(X_train, y_train)

# Cross-validation
cv_scores = cross_val_score(rf_model, X_train, y_train, cv=5, scoring='r2')
print(f"Cross-validation R2: {cv_scores.mean():.3f} (+/- {cv_scores.std()*2:.3f})")

# Predictions
y_pred = rf_model.predict(X_test)

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
mae = mean_absolute_error(y_test, y_pred)
mape = np.mean(np.abs((y_test - y_pred) / y_test)) * 100
r2 = r2_score(y_test, y_pred)

print("\n==== Model Evaluation ===")
print(f"RMSE: {rmse:.2f}")
print(f"MAE : {mae:.2f}")
print(f"MAPE: {mape:.2f}%")
print(f"R2 : {r2:.3f}")

# Plot Actual vs Predicted and Feature Importance
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.scatter(y_test, y_pred, alpha=0.6, s=20)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.xlabel('Actual Price ($)')
plt.ylabel('Predicted Price ($)')
plt.title(f'Actual vs Predicted Price (R2 = {r2:.3f})')
plt.grid(alpha=0.3)

importance_df = pd.DataFrame({
    'Feature': core_features,
    'Importance': rf_model.feature_importances_
}).sort_values('Importance', ascending=False)

plt.subplot(1, 2, 2)
sns.barplot(x='Importance', y='Feature', data=importance_df.head(8))
plt.title('Top 8 Most Important Features')
plt.tight_layout()
plt.show()

# Error analysis
abs_errors = np.abs(y_test - y_pred)
print(f"\nMean Absolute Error: ${abs_errors.mean():.2f}")
print(f"Median Absolute Error: ${np.median(abs_errors):.2f}")

price_bins = pd.cut(y_test, bins=5)
error_by_price = pd.DataFrame({
    'Absolute_Error': abs_errors,
    'Actual': y_test
}).groupby(price_bins).agg({'Absolute_Error': ['mean', 'median'], 'Actual': 'count'}).round(2)

print("\nError Analysis by Price Range:")
print(error_by_price)

```

选择文件 未选择文件

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving data.csv to data.csv

Original number of rows: 4600

Original columns: ['date', 'price', 'bedrooms', 'bathrooms', 'sqft\_living', 'sqft\_lot', 'floors', 'waterfront', 'view', 'condition', 'sqft\_above', 'sqft

Number of rows after cleaning: 3757

	price	sqft_living	bedrooms	bathrooms
count	3.757000e+03	3757.000000	3757.000000	3757.000000
mean	4.731519e+05	1926.822199	3.316210	2.049241
min	7.800000e+03	370.000000	1.000000	0.750000
25%	3.100000e+05	1390.000000	3.000000	1.500000
50%	4.370000e+05	1850.000000	3.000000	2.000000
75%	5.973260e+05	2380.000000	4.000000	2.500000
max	1.150000e+06	4130.000000	9.000000	5.750000
std	2.074133e+05	706.241941	0.869606	0.690970

Training set: 2629 records, Test set: 1128 records

Cross-validation R<sup>2</sup>: 0.542 (+/- 0.046)

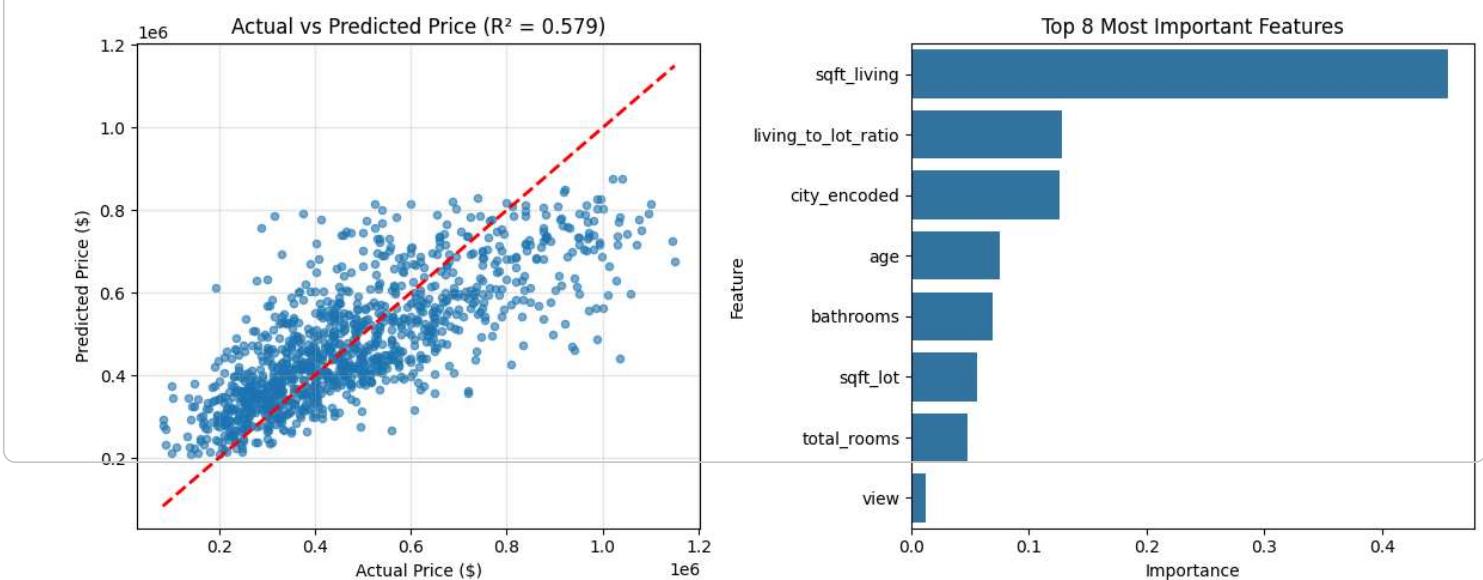
==== Model Evaluation ====

RMSE: \$138,191.03

MAE : \$106,776.96

MAPE: 26.04%

R<sup>2</sup> : 0.579



Mean Absolute Error: \$106,776.96

Median Absolute Error: \$86,186.04

Error Analysis by Price Range:

price	Absolute_Error	Actual		
		mean	median	count
(81933.0, 296400.0]	104397.13	91525.97	227	
(296400.0, 509800.0]	78355.48	58088.16	480	
(509800.0, 723200.0]	103070.78	94685.05	259	
(723200.0, 936600.0]	157987.02	148532.73	112	
(936600.0, 1150000.0]	294915.05	275718.31	50	